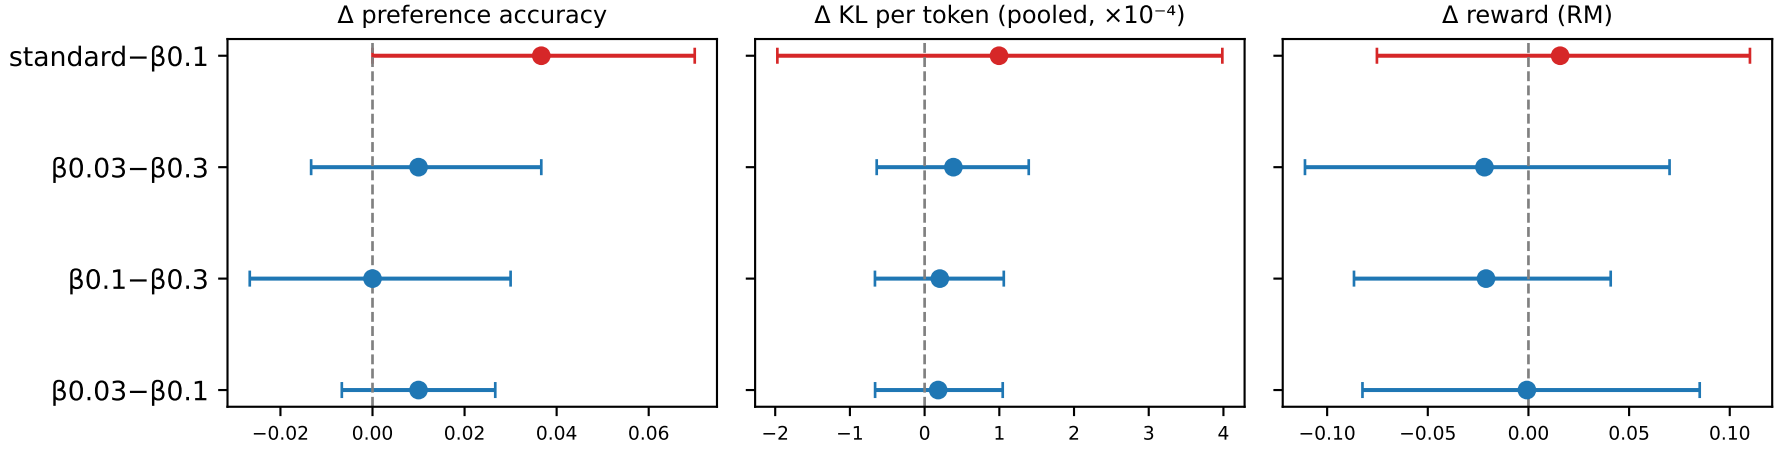
Δ preference accuracy
Δ KL per token (pooled, ×10⁻⁴)
Δ reward (RM)
standard−β0.1
β0.03−β0.3
β0.1−β0.3
β0.03−β0.1
−0.02
0.00
0.02
0.04
0.06
−2
−1
0
1
2
3
4
−0.10
−0.05
0.00
0.05
0.10