




learned reward



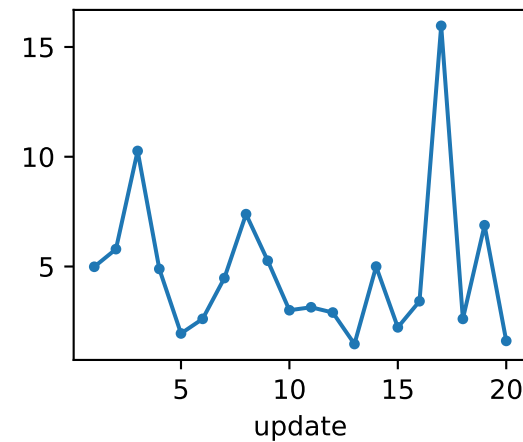
KL to ref (token)



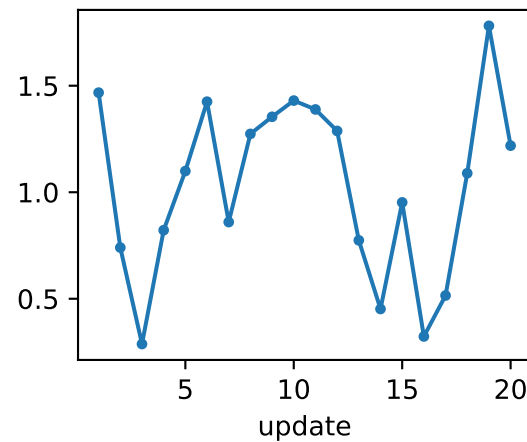
policy loss



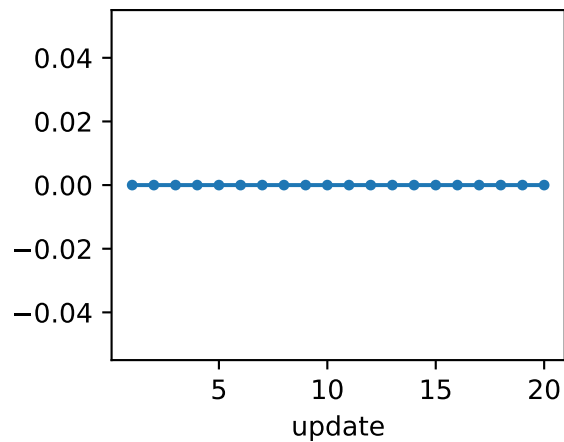
value loss



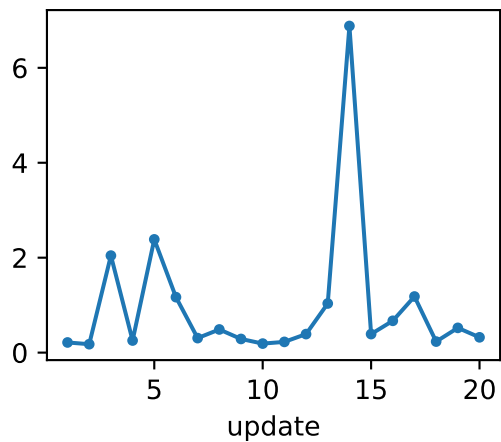
entropy



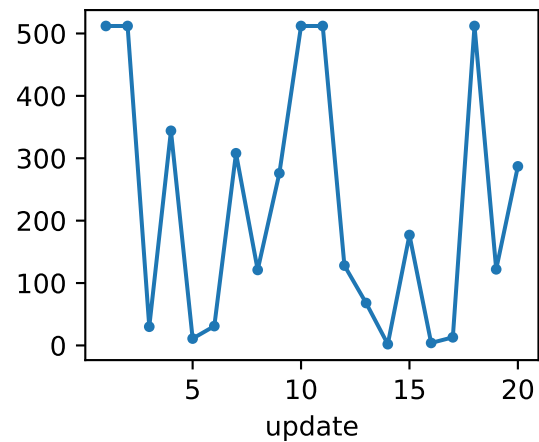
clip fraction



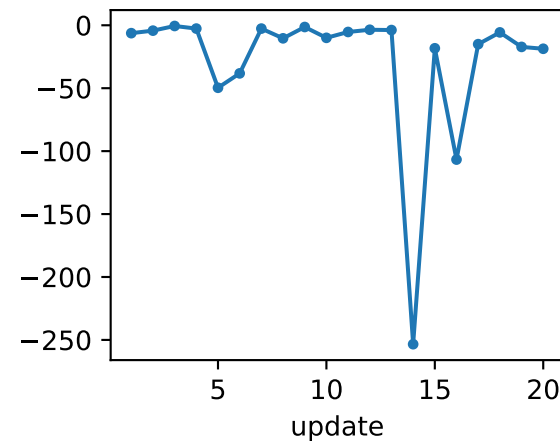
policy grad norm



response tokens



critic explained var.



policy change per update (k3)

